


# Project 3 of Robotics Nanodegree Object Detection



October 14, 2017

In this project, we use point cloud data (from topic `/pr2/world/points`) and a perception pipeline to detect and identify objects on a table. The implementation of the pipeline involves multiple steps that are outlined in the sections below. The pipeline is used to identify a set of 3, 5 and 8 objects.

## 0.1 Perception pipeline

### 0.1.1 Downsampling

The first step consists in using a Voxel Grid Filter. The voxel size (`LEAF_SIZE`) is set to 0.01 m. A higher number results in a loss of resolution.

```
cloud = ros_to_pcl(pcl_msg)
vox = cloud.make_voxel_grid_filter()
LEAF_SIZE = 0.01
vox.set_leaf_size(LEAF_SIZE, LEAF_SIZE, LEAF_SIZE)
cloud_filtered = vox.filter()
```

### 0.1.2 Filter noise

Next, we use a statistical outlier filter to *clean up* the point cloud: *i.e* remove the noise. For each data point, we compute the distance from that point to all the  $k$  neighboring points. Assuming a Gaussian distribution, the mean  $\mu$  and standard deviation  $\sigma$  are computed. Any points that are outside the range  $\mu \pm x \times \sigma$  are labeled as outliers and removed from the point cloud.  $x$  is a scaling factor of the standard deviation..




Figure 1: Top image is the original point cloud, and bottom image is the point cloud after statistical outlier filtering.

In our implementation, we set the number of neighboring points to be  $k = 50$ , and the scaling factor  $x = 0.1$ .

```
outlier_filter = cloud_filtered.make_statistical_outlier_filter()
```

```

outlier_filter.set_mean_k(50)
x = 0.1
outlier_filter.set_std_dev_mul_thresh(x)
cloud_filtered_outliers = outlier_filter.filter()

```

### 0.1.3 Passthrough filter

The passthrough filter is a crop-like method that enables to reduce the field of view. The first passthrough filter is along the  $z$  axis, where only the points with  $z$  coordinates in the range  $[0.6, 1.1]$  are preserved.

```

passthrough = cloud_filtered_outliers.make_passthrough_filter()
filter_axis = 'z'
passthrough.set_filter_field_name(filter_axis)
axis_min = 0.6
axis_max = 1.1
passthrough.set_filter_limits(axis_min, axis_max)
cloud_filtered = passthrough.filter()
passthrough = cloud_filtered.make_passthrough_filter()

```

However, a single passthrough filter is not sufficient. In fact, the front edges of the 2 drop-boxes remains in the view point of the camera and can alter the object recognition. Therefore, in addition of the  $z$  axis passthrough, we also use a passthrough filter along the  $y$  axis in the range  $[-0.35, 0.35]$ .

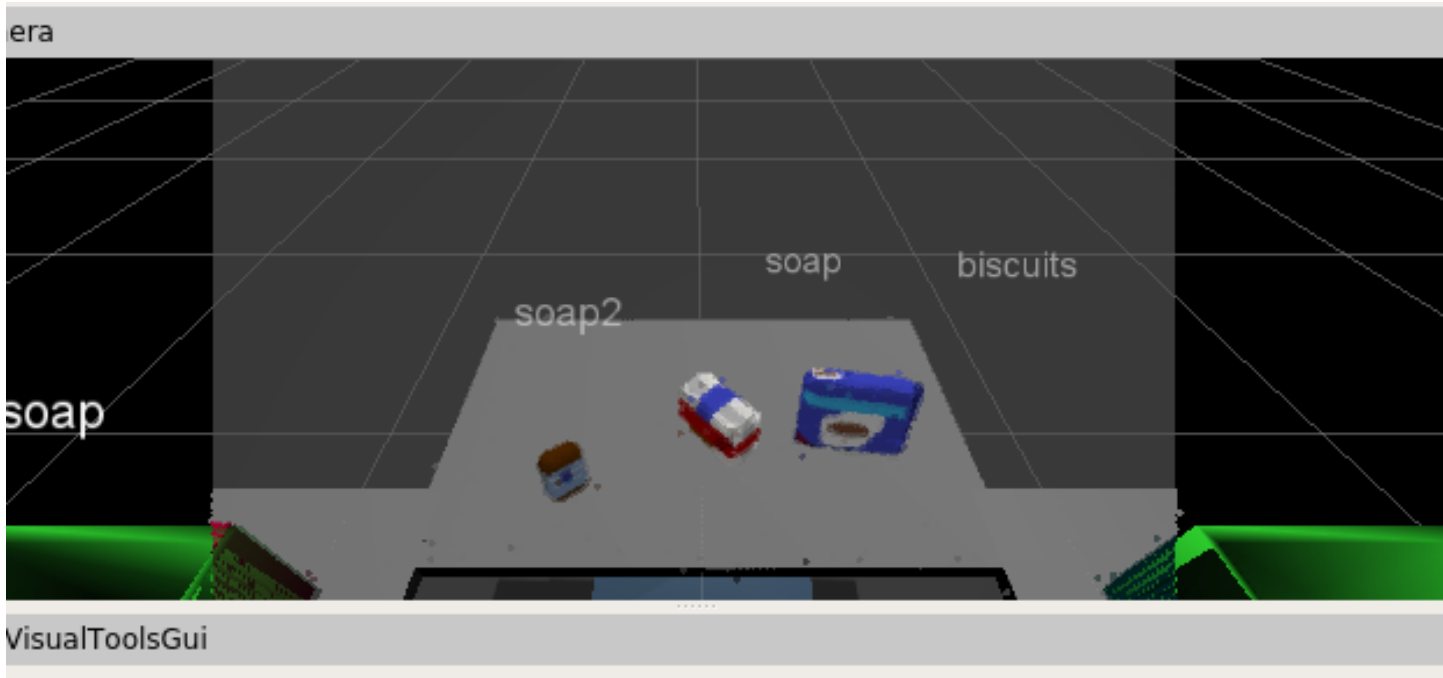


Figure 2: With a passthrough filter along  $z$  direction, there remains point clouds because related to the edges of the box (left and right). Those points can lead to erroneous recognition.

```


passthrough = cloud_filtered.make_passthrough_filter()
filter_axis = 'y'
passthrough.set_filter_field_name(filter_axis)

```

```
axis_min = -0.35
axis_max = 0.35
passthrough.set_filter_limits(axis_min, axis_max)
cloud_filtered = passthrough.filter()
```

### 0.1.4 RANSAC

In this step, we use RANSAC (Random Sample Consensus) algorithm to separate the objects and the table. The table is modeled as a plane to which we fit the entire dataset. We can then extract 2 datasets: the inliers and the outliers. The inliers are the data points that follow the models and as a consequence are most likely associated with the table.




```
seg = cloud_filtered.make_segmenter()
seg.set_model_type(pcl.SACMODEL_PLANE)
seg.set_method_type(pcl.SAC_RANSAC)
max_distance = 0.01
seg.set_distance_threshold(max_distance)
inliers, coefficients = seg.segment()
cloud_table = cloud_filtered.extract(inliers, negative=False)
cloud_objects = cloud_filtered.extract(inliers, negative=True)
```

The hyper-parameter is the *max\_distance*: it represents the maximum distance that the datapoint can be from the model best fit and still be considered to be an inlier. We set *max\_distance*= 0.01.

### 0.1.5 Clustering (DBSCAN)

The *DBSCAN* (Density Based Spatial Clustering of Applications with Noise) is a unsupervised learning algorithm used for segmentation. The 2 hyper-parameters are *min\_samples*, the minimum number of points that comprise a cluster, and *max\_samples* is the maximum distance between cluster points.



```

white_cloud = XYZRGB_to_XYZ(cloud_objects)
tree = white_cloud.make_kdtree()
ec = white_cloud.make_EuclideanClusterExtraction()
ec.set_ClusterTolerance(0.05)
ec.set_MinClusterSize(100)
ec.set_MaxClusterSize(1550)
ec.set_SearchMethod(tree)
cluster_indices = ec.Extract()
cluster_color = get_color_list(len(cluster_indices))
color_cluster_point_list = []

for j, indices in enumerate(cluster_indices):
    for i, indice in enumerate(indices):
        color_cluster_point_list.append([white_cloud[indice][0],
                                         white_cloud[indice][1],
                                         white_cloud[indice][2],
                                         rgb_to_float(cluster_color[j])])

cluster_cloud = pcl.PointCloud_PointXYZRGB()
cluster_cloud.from_list(color_cluster_point_list)

```

We set the tuning parameters as follows: *cluster\_tolerance*= 0.05, *MinClusterSize*= 100, *MaxClusterSize*= 1550.

### 0.1.6 Object Recognition

Now that we have segmented the point cloud, we need an algorithm to recognize the objects/segments. First, we need to define a set of features to characterize the objects, and then use a supervised learning algorithm to classify the clusters.

#### Create features

The features are composed of the color histogram and the normal surface histogram. The color histogram is constructed from the HSV channels, with the pixel intensity range  $[0, 255]$  splitted into 16 bins. The normal surface histogram has a range  $[0, 512]$  and 32 bins. The 2 histograms are then concatenated into a single vector and normalized.

```
chists = compute_color_histograms(sample_cloud, using_hsv=True)
normals = get_normals(sample_cloud)
nhists = compute_normal_histograms(normals)
feature = np.concatenate((chists, nhists))
```

#### Train model

The classification algorithm is a Support Vector Machine. To train the model, we create a balanced dataset of different objects: *soap*, *soap2* and *biscuits* for example. We trained our model using features from 150 samples per category. The decision boundary can be either linear or non-linear, so the choice of the kernel in the *SVM* affects the performance of the classifier. We ran a Grid search to find the best performing model:

```
parameter_grid = [{ 'C':[0.1, 1, 10, 100, 1000], 'kernel': [ 'linear' ] },
                   { 'C': [0.1, 1, 10, 100, 1000], 'gamma': [0.01, 0.001, 0.0001] },
                   { 'C': [0.1, 1, 10, 100, 1000], 'gamma': [0.1, 0.01, 0.001, 0.0001] },
clf = GridSearchCV(estimator=svm.SVC(), param_grid=parameter_grid, n_jobs=-1)
....
....
clf.fit(X=X_train, y=y_train)
print('Best_C:', clf.best_estimator_.C)
print('Best_Kernel:', clf.best_estimator_.kernel)
print('Best_Gamma:', clf.best_estimator_.gamma)
```

We also use a  $K - fold$  cross validation to train the model. Such an approach is particularly useful when dealing with a small dataset.

The best model is with  $kernel=linear$ ,  $C = 1$ .

#### Results

We trained 3 models depending on the number of objects to be detected.




Figure 3: Normalized Confusion matrix for 3 objects (test1). The model has an accuracy of 96.89%

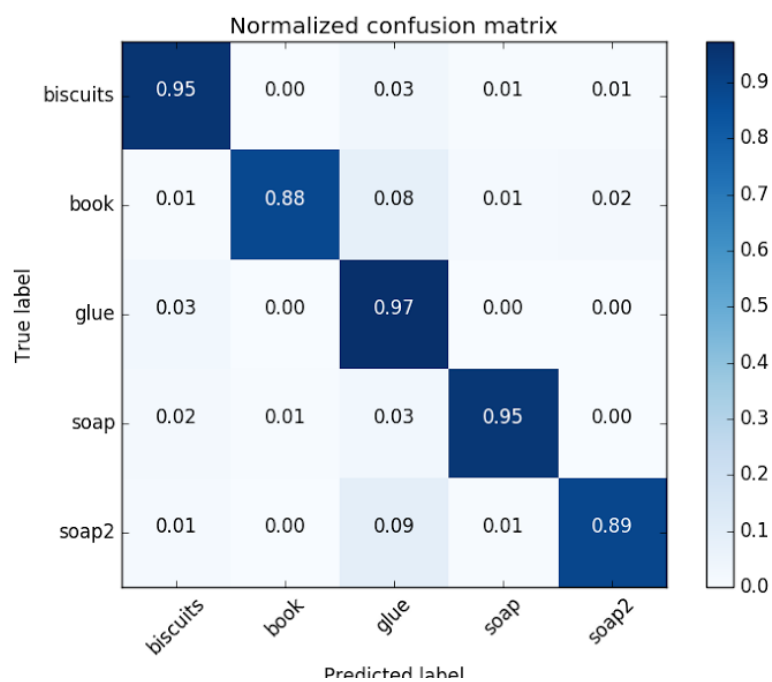


Figure 4: Normalized Confusion matrix for 5 objects (test2). The model has an accuracy of 92.78%




Figure 5: Normalized Confusion matrix for 8 objects (test3). The model has an accuracy of 93%

## 0.2 Pipeline in action

Below we show the output of the pipeline for the 3 worlds, where 3, 5 and 8 objects must be identified.

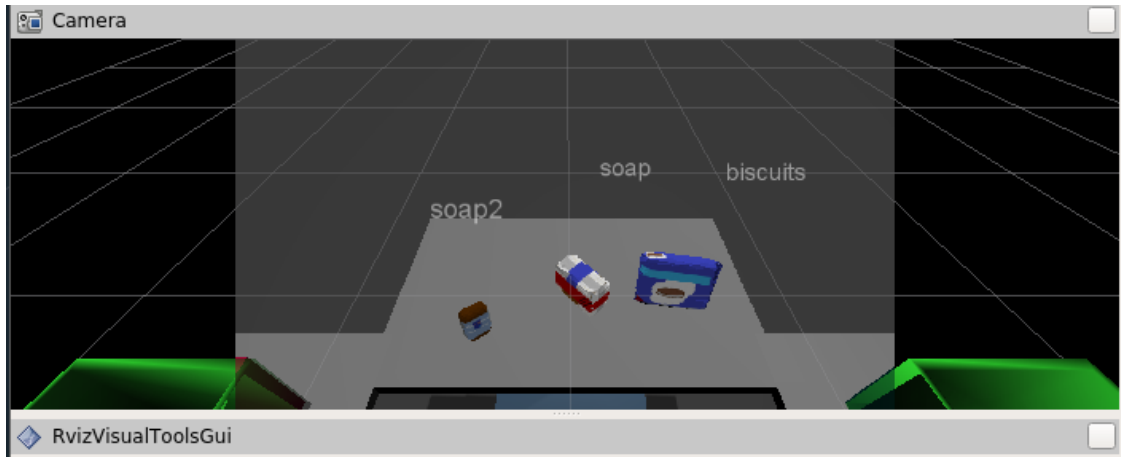


Figure 6: 3 out of 3 objects correctly identified





Figure 7: 5 out of 5 objects correctly identified



Figure 8: 6 out of 8 objects correctly identified. The model cannot separate the glue and the book, and merged into a single cluster that it labels erroneously as soap.